



# Neo4j Database Management System

Akshay Pudage  
Kushal Gevaria



## Problems with other DBMS

- Traditional relational DBMS cannot handle unstructured data and ad-hoc relations
- Relations in RDBMS need expensive join operations
- NoSQL databases are bad at handling relations
- Reciprocal queries can be very expensive

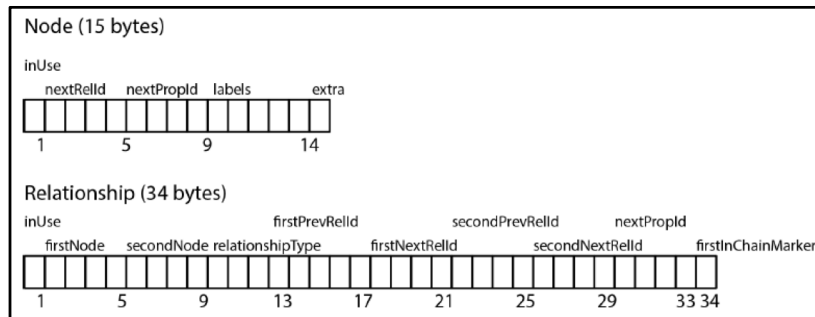


# Why Neo4j?

- Uses graph data structure for modeling data
- Relations can be easily represented in graphs
- Graph traversals take  $O(1)$  time

# Data Storage & Indexing

- Uses index-free adjacency which speeds up traversals
- Native graph processing capability
- Maintain separate stores for storing nodes, relationships and properties.
- Uses fixed record size to enable  $O(1)$  lookups
- Forms a linked list internally for navigating across records





# Query Processing and Optimization

- Cypher: The most declarative language
- CRUD operations like create, read, update and delete operations can be done on the database using simple Cypher queries
- Cypher query language represents all the complex relations and structures in a meaningful and concise way
- Cypher provides a special operator to find the neighborhood nodes the way we want. In traditional relational database, this is cumbersome with multiple joins and cartesian products.



# Query Processing and Optimization

- Regular search over the entire data can be costly in terms of time and resources
- Cypher provides an option of indexing. We can create an index on a single attribute or composite indexes are also available
- Composite indexes allow index creation on more than one attribute
- The Neo4j tries to execute the queries as fast as possible and in order to do that sometimes a query tuning
- Each cypher query gets optimized and transformed into an execution plan.
- Execution plan: Uses minimal resources chosen from this set of plans.
- For parameterized queries as oppose to hardcoded literals the query engine also re-uses the same execution plan instead of building all plans and choosing the best.
- Furthermore, there are several operators that aid in choosing the best query plan.



# Transaction Management & Security

- Uses two-commit transactions
- Lock manager acquires write locks on nodes and relationships
- If successful, changes are flushed to disk and locks are released
- If failure, changes are discarded and locks are released
- Creates a transaction object internally which keeps track of changes
- Uses Write Ahead Log while committing changes to disk



# Neo4j NoSQL Database Application

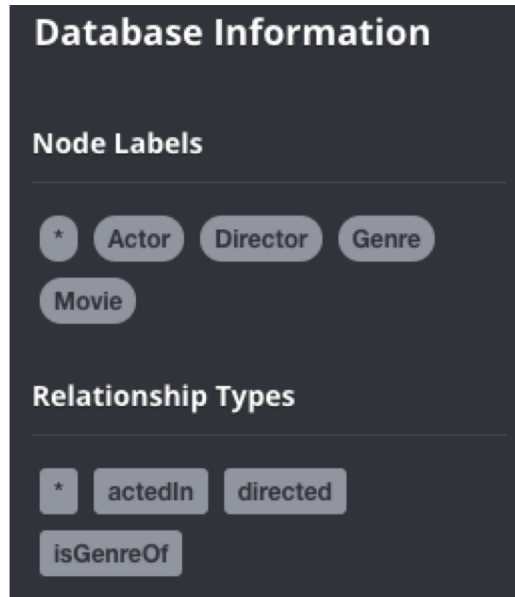
- Neo4j community version 3.2.5: Application UI that operates conveniently across all the platforms
- Backend is supported by Java: Neo4j works best with Java and is supported by all the packages in Java
- “Org.neo4j.graphdb”: To establish the connection with the database, using "GraphDatabaseService" class
- “GraphDatabaseService” package provides an instance of Neo4j graphDB to perform CRUD operations on the nodes and relationships in java





# Database Information

- Nodes are created with their respective labels
- Movie, Actor, Director, Genre are the type of nodes. Here, type is displayed by the labels assigned to those nodes
- Relations are created between nodes. Each relation has a relation type. It can be directed or undirected.
- For IMDB dataset, we created three basic relations between nodes.
- “ActedIn”: Relation between an actor and a movie
- “Directed”: Relation between a director and a movie
- “IsGenreOf”: Relation between a genre and a movie



# Neo4j Sample Graph for IMDB

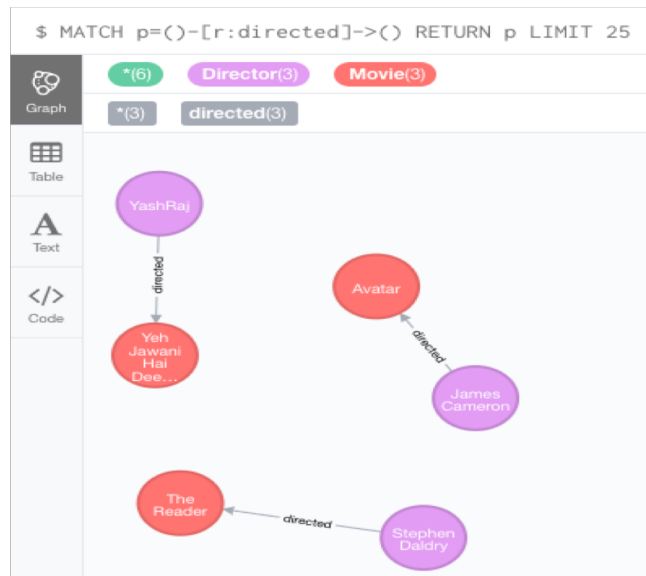
- Small subgraph of the IMDB dataset after certain CRUD operations is displayed on the right
- There are currently 4 actors, 3 directors, 4 genres and 3 movies in this sample graph.
- Whole graph has millions of nodes with complex relationships between multiple nodes.
- Movie “Avatar” has two genres namely “Animated”, “Drama”, one director “James Cameron” and an actress “Zoe Saldana”. (In reality there is more to this graph)



# Relationship between Directors and Movies

- On the right is the example of Cypher query to get all the relations between directors and movies
- Similar can be achieved in Java ->

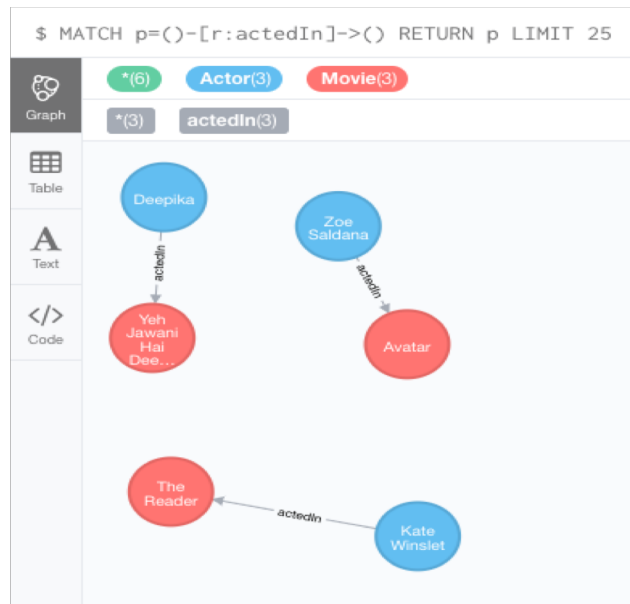
```
Node dNode = getDbSiIMDB().findNode(directorLabel,
"name", directorName);
for(Relationship r : dNode.getRelationships()) {
    // get nodes on other end of this relationship
    Node tempMovie = r.getEndNode();
}
```



# Relationship between Actors and Movies

- On the right is the example of Cypher query to get all the relations between actors and movies
- Similar can be achieved in Java ->

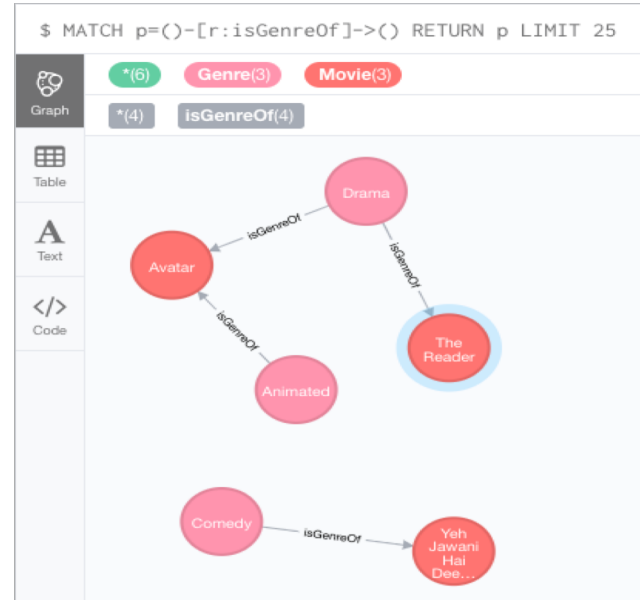
```
Node aNode = getDbSiIMDB().findNode(actorLabel,
"name", actorName);
for(Relationship r : aNode.getRelationships()) {
    // get nodes on other end of this relationship
    Node tempMovie = r.getEndNode();
}
```



# Relationship between Genres and Movies

- On the right is the example of Cypher query to get all the relations between genres and movies
- Similar can be achieved in Java ->

```
Node gNode = getDbSilMDB().findNode(genreLabel,
"name", genreName);
for(Relationship r : gNode.getRelationships()) {
    // get nodes on other end of this relationship
    Node tempMovie = r.getEndNode();
}
```





# Create Nodes and Relationships using Java

Example to create a director node:

```
Node dNode = getDbSiIMDB().createNode();  
dNode.setProperty("name", director.getDirector());
```

Example to create a label for all director nodes:

```
Label directorLabel = Label.label("Director"); // label  
of director  
dNode.addLabel(directorLabel);
```

Example to create a relationship type for all director nodes:

```
RelationshipType directorMovieConnection =  
    RelationshipType.withName("directed");  
dNode.createRelationshipTo(mNode,  
    directorMovieConnection);
```



# Find Nodes and Relationships using Java

Example to find a director node  
and its relationships::

```
Node dNode = getDbSiIMDB().findNode(directorLabel,  
    "name", directorName);  
for(Relationship r : dNode.getRelationships()) {  
    // get nodes on other end of this relationship  
    Node tempMovie = r.getEndNode();  
}
```



# Update Nodes using Java

Example to update a director node:

```
Node dNode = getDbsiIMDB().findNode(directorLabel,  
    "name", oldName);  
dNode.setProperty("name", newName);
```





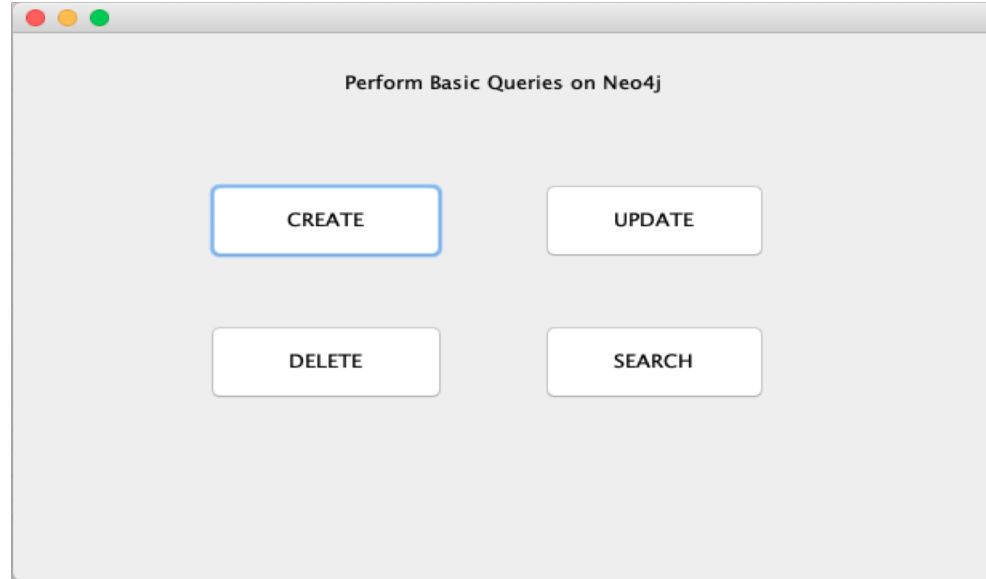
# Delete Nodes and its Relationships using Java

Example to delete a director node and its relationships:

```
Node dNode = getDbsiIMDB().findNode(directorLabel,
    "name", name);
for(Relationship r : dNode.getRelationships()) {
    r.delete();
}
dNode.delete();
```

# Neo4j GUI Application using Java Spring Builder

CRUD Operation Panel



# Neo4j GUI Application: CREATE MOVIE

Create a movie



The screenshot shows a web application window titled "CREATE" with a standard macOS-style title bar (red, yellow, green buttons). In the top left corner, there is a button labeled "<- Back". The main content area contains four input fields, each with a label to its left: "Movie:", "Genre:", "Director:", and "Actor:". The "Movie:" field contains the text "Avengers: Infinity V" and is highlighted with a blue border. The "Genre:" field contains "Fantasy", the "Director:" field contains "Joe Russo", and the "Actor:" field contains "Chris Hemsworth". To the right of these fields, there is a status message: "Status: Successfully added the movie" in blue text. At the bottom right of the form, there is a button labeled "Create".

# Neo4j GUI Application: SEARCH BY DIRECTOR

Search Movies by Director

SEARCH

<- Back

Search Movie by Director

Director:

Search Movie by Actor

Actor:

Search Movie by Genre

Genre:

Movies

| Movies     |
|------------|
| The Reader |

Status: Found some movies for director: Stephen Daldry

# Neo4j GUI Application: SEARCH BY ACTOR

Search Movie by Actor

The screenshot shows a web application window titled "SEARCH". It has three search sections: "Search Movie by Director", "Search Movie by Actor", and "Search Movie by Genre". Each section has a text input field and a "Search by [Category]" button. The "Search Movie by Actor" section has "Kate Winslet" entered in the input field, which is highlighted with a blue border. To the right of the search sections is a "Movies" list. The list has a header "Movies" and one item "The Reader". At the bottom of the window, there is a "Status:" label followed by the text "Found some movies for actor: Kate Winslet" in blue.

| SEARCH   |   |
|--|---|
| <a href="#">&lt;- Back</a>                               |   |
| <b>Search Movie by Director</b>                          |   |
| Director:  | <input type="text"/> <a href="#">Search by Director</a>                   |
| <b>Search Movie by Actor</b>                             |   |
| Actor:   | <input type="text" value="Kate Winslet"/> <a href="#">Search by Actor</a> |
| <b>Search Movie by Genre</b>                             |   |
| Genre:   | <input type="text"/> <a href="#">Search by Genre</a>                      |
| <b>Status:</b> Found some movies for actor: Kate Winslet |   |

**Movies**

| Movies     |
|------------|
| The Reader |

# Neo4j GUI Application: SEARCH BY GENRE

Search Movie by Genre

The screenshot shows a web application interface for searching movies. It has a title bar with red, yellow, and green window control buttons. The main heading is "SEARCH". Below it is a "<- Back" button. There are three search sections: "Search Movie by Director" with a "Director:" label and a "Search by Director" button; "Search Movie by Actor" with an "Actor:" label and a "Search by Actor" button; and "Search Movie by Genre" with a "Genre:" label, a text input field containing "Fantasy", and a "Search by Genre" button. To the right of these sections is a "Movies" list box. The list box has a title "Movies" and contains two entries: "The Avengers" and "Avengers: Infinity Wars". At the bottom of the interface, there is a "Status:" label followed by the text "Found some movies for genre: Fantasy" in blue.

| Search Method            | Field Label | Field Value | Action Button      |
|--------------------------|-------------|-------------|--------------------|
| Search Movie by Director | Director:   |             | Search by Director |
| Search Movie by Actor    | Actor:      |             | Search by Actor    |
| Search Movie by Genre    | Genre:      | Fantasy     | Search by Genre    |

Status: Found some movies for genre: Fantasy

Movies

- The Avengers
- Avengers: Infinity Wars

# Neo4j GUI Application: UPDATE GENRE

Update Misspelled Genre

<- Back

UPDATE

|           | From                               | To                                  |                 |
|-----------|------------------------------------|-------------------------------------|-----------------|
| Movie:    | <input type="text"/>               | <input type="text"/>                | Update Movie    |
| Genre:    | <input type="text" value="Actin"/> | <input type="text" value="Action"/> | Update Genre    |
| Director: | <input type="text"/>               | <input type="text"/>                | Update Director |
| Actor:    | <input type="text"/>               | <input type="text"/>                | Update Actor    |

Status: Genre Updated: Action

# Neo4j GUI Application: UPDATE MOVIE

Update Misspelled Movie

<- Back UPDATE

|           | From                 | To                  |                 |
|-----------|----------------------|---------------------|-----------------|
| Movie:    | ngers: Infinity Wars | ngers: Infinity War | Update Movie    |
| Genre:    |                      |                     | Update Genre    |
| Director: |                      |                     | Update Director |
| Actor:    |                      |                     | Update Actor    |

Status: Movie Updated: Avengers: Infinity War



# Neo4j GUI Application: DELETE DIRECTOR

Delete a Director

<- Back DELETE

|           |  |                 |
|-----------|--|-----------------|
| Movie:    | <input type="text"/>                   | Delete Movie    |
| Genre:    | <input type="text"/>                   | Delete Genre    |
| Director: | <input type="text" value="Joe Russo"/> | Delete Director |
| Actor:    | <input type="text"/>                   | Delete Actor    |

Status: Successfully deleted director with it's relationships: Joe Russo



## References

[1] E. Eifrem. Graph database service, 2002.

<https://neo4j.com/docs/java-reference/current/javadocs/org/neo4j/graphdb/GraphDatabaseService.html>

[2] E. Eifrem. Neo4j cypher refcard 3.4, 2002.

<https://neo4j.com/docs/cypher-refcard/current/>

[3] E. Eifrem. Neo4j cypher query language, 2002

<https://neo4j.com/developer/cypher-query-language/>

[4] E. Eifrem. Neo4j cypher query tuning, 2002

<https://neo4j.com/docs/developer-manual/current/cypher/query-tuning/>

[5] Neo4j. Chapter 7: Security. Accessed: 2018-10-26.

[6] I. Robinson, J. Webber, and E. Eifrem. Graph databases. O'Reilly Media, Inc., 2013.



# Thank You